

DB Programming

John Strickler

Version 1.0, December 2025

Table of Contents

Chapter 1: Database Access	1
The DB API	2
Connecting to a Server	4
<code>connect()</code> examples	5
Creating a Cursor	7
Querying data	8
Fetch methods	11
Non-query statements	14
SQL Injection	17
Parameterized Statements	19
Metadata	31
Dictionary Cursors	34
Generic alternate cursors	38
Transactions	42
Autocommit	44
Object-relational Mappers	45
NoSQL	46
Index	49

Chapter 1: Database Access

Objectives

- Understand the Python DB API architecture
- Connect to a database
- Execute simple and parameterized queries
- Fetch single and multiple row results
- Execute non-query statements
- Get metadata about a query
- Start transactions and commit or rollback as needed

The DB API

- Most popular Python DB interface
- Specification, not abstract class
- Many modules for different DB implementations
- Hides actual DBMS implementation

To make database programming simpler and more consistent, Python provides the DB API. This is an API to standardize working with databases. When a package is written to access a database, it is written to conform to the API, and thus programmers do not have to learn a new set of methods and functions for each different database architecture.

DB API objects and methods

```
conn = package.connect(connection-arguments)
cursor = conn.cursor()
num_lines = cursor.execute(query)
num_lines = cursor.execute(query-with-placeholders, iterable)
num_lines = cursor.executemany(query-with-placeholders, iterable)
all_rows = cursor.fetchall()
some_rows = cursor.fetchmany(n)
one_row = cursor.fetchone()
conn.commit()
conn.rollback()
```

Table 1. Available Interfaces (using Python DB API-2.0)

Database	Python package
Firebird (and Interbase)	KInterbasDB
Databricks	databricks-sql-connector
Elasticsearch	elasticsearch-dbapi
Excel	excel-dbapi
Google BigQuery	google-cloud-bigquery
IBM DB2	ibm-db
Informix	informixdb
Ingres	ingmod
Microsoft SQL Server	pymssql
MySQL	pymysql
ODBC	pyodbc
Oracle	oracledb cx_oracle (obsolete as of 2022)
PostgreSQL	psycopg (previously psycopg2)
SAP DB (also known as "MaxDB")	sapdbapi
SQLite	sqlite3
Sybase	Sybase



This list is not comprehensive, and there are additional interfaces available for some of the listed DBMSs.

Connecting to a Server

- Import appropriate library
- Use connect() to get a connection object
- Specify host, database, username, password, etc.

To connect to a database server, import the package for the specific database. Use the package's `connect()` method to get a connection object. The `connect()` function requires the information needed to access the database, which may include the host, initial database, username, or password.

Argument names for the `connect()` method are not consistent across packages. Most `connect()` methods use individual arguments, such as `host`, `database`, etc., but some use a single string argument.

When finished with the connection, call the `close()` method on the connection object. Many database modules support the context manager `with` statement, and will automatically close the database when the `with` block is exited. Check the documentation to see how this is implemented for a particular database.

Example

```
import pymysql

conn = pymysql.connect (host = "dbserver",
                       user = "adeveloper",
                       passwd = "s3cr3t",
                       db = "samples")
# Interact with database here ...
conn.close()
```

```
import sqlite3

with sqlite3.connect('sample.db') as conn:
    # Interact with database here ...
```

connect() examples

PostgreSQL using `pyscopg`

```
conn = psycopg.connect(  
    host="localhost",  
    dbname="postgres",  
    user="adeveloper",  
    password='$3cr3t',  
)
```

Oracle using `oracledb`

```
conn = oracledb.connect(  
    user='adeveloper',  
    password='$3cr3t',  
    dsn="dbhost.example.com/orclpdb"  
)
```

```
conn = oracledb.connect(  
    user='adeveloper',  
    password='$3cr3t',  
    host="localhost",  
    port=1521,  
    servicename="orclpdb",  
)
```

SQLite3 using `sqlite3`

```
conn = sqlite3.connect('testdb') # on-disk database(single file)
```

```
conn = sqlite3.connect(':memory:') # in-memory database
```

Microsoft SQLServer using `pymssql`

```
conn = pymssql.connect(  
    server=r"HOST\INSTANCE",  
    user="adeveloper", # SQLServer user, not Windows user  
    password="s3cr3t", # SQLServer password  
    database="testdb")
```

```
conn = pymysql.connect (
    dsn="DSN",
)
```

MySQL/MariaDB using `pymysql`

```
conn = pymysql.connect (
    host="localhost",
    user="adeveloper",
    passwd="$3cr3t",
    db="testdb",
)
```

Any ODBC-compliant DB using `pyodbc`

```
conn = pyodbc.connect('''
    DRIVER={SQL Server};
    SERVER=localhost;
    DATABASE=testdb;
    UID=adeveloper;
    PWD=$3cr3t
''')
```

```
conn = pyodbc.connect('DSN=testdsn;PWD=$3cr3t')
```



`pyodbc.connect()` has one (string) parameter, not multiple parameters

IBM DB2 using `ibm-db`

```
import ibm_db_dbi as db2
conn = db2.connect(
    "DATABASE=testdb;HOSTNAME=localhost;PORT=50000;PROTOCOL=TCPIP;UID=db2inst1;PWD=scripts;",
    "",
    ""
)
```

Creating a Cursor

- Cursor can execute SQL statements
- Create with `cursor()` method
- Multiple cursors available
 - Standard cursor
 - Returns rows as tuples
 - Other cursors
 - Return dictionary
 - Return hybrid dictionary/list
 - Leave data on server

Once you have a connection object, call `cursor()` to create a cursor object. A cursor is an object that can execute SQL code and fetch results. Each connection may have one or more active cursors.

The default cursor for most packages returns each row as a tuple of values. There are optional cursors that can return data in different formats, or that control whether data is stored on the client or the server.



The examples in this chapter are implemented with SQLite, since the `sqlite3` module is part of the standard library. Most of the examples are also implemented for PostgreSQL and MySQL. See `db_mysql_*.py` and `db_postgres_*.py` in EXAMPLES.

Querying data

- **`cursor.execute(query)`**
 - Gets all data from query
 - Returns # rows in result set
- Cursor is iterator over result set
- Return rows as tuples of values

Once you have a cursor, you can use it to execute queries via the `execute()` method. The first argument to `execute()` is a string containing one SQL statement.

For queries, `execute()` returns the number of rows in the result set. For non-query statements, `execute()` returns the number of rows affected by the operation.

The cursor object is an iterator over the rows in the result set.

For standard cursors, all data is transferred from the database server to your program's memory when `execute()` is called.



For `sqlite3`, `execute()` returns the cursor object, so you can say `execute(QUERY-STATEMENT).fetchall()`.

Example

db_sqlite_basics.py

```
import sqlite3

# conn = sqlite3.Connection(...)
with sqlite3.connect("../DATA/presidents.db") as conn: # connect to the database

    s3_cursor = conn.cursor() # get a cursor object

    # select specified columns from all presidents
    s3_cursor.execute('''
        select termnum, firstname, lastname, party
        from presidents
    ''') # execute a SQL statement

    for term, firstname, lastname, party in s3_cursor:
        print(f"{term:2d} {firstname:25} {lastname:20} {party}")
print()
```

db_sqlite_basics.py

1 George	Washington	no party
2 John	Adams	Federalist
3 Thomas	Jefferson	Democratic - Republican
4 James	Madison	Democratic - Republican
5 James	Monroe	Democratic - Republican
6 John Quincy	Adams	Democratic - Republican
7 Andrew	Jackson	Democratic
8 Martin	Van Buren	Democratic
9 William Henry	Harrison	Whig
10 John	Tyler	Whig

...

39 James Earl 'Jimmy'	Carter	Democratic
40 Ronald Wilson	Reagan	Republican
41 George Herbert Walker	Bush	Republican
42 William Jefferson 'Bill'	Clinton	Democratic
43 George Walker	Bush	Republican
44 Barack Hussein	Obama	Democratic
45 Donald John	Trump	Republican
46 Joseph Robinette	Biden	Democratic
47 Donald John	Trump	Republican

Fetch methods

- Alternate way to get rows
 - `fetchone()` get next row
 - `fetchmany()` get N rows
 - `fetchall()` get all remaining rows

Cursors provide three methods for returning query results. These can be used instead of iterating over the cursor.

`fetchone()` returns the next available row from the query results.

`fetchmany(n)` returns up to n rows. This is useful when the query returns a large number of rows.

`fetchall()` returns a tuple of all rows.

For all three methods, each row is returned as a tuple of values.

Example

db_sqlite_fetch.py

```
import sqlite3

# conn = sqlite3.Connection(...)
with sqlite3.connect("../DATA/presidents.db") as conn: # connect to the database

    s3_cursor = conn.cursor() # get a cursor object

    # select specified columns from all presidents
    s3_cursor.execute('''
        select termnum, firstname, lastname, party
        from presidents
    ''') # execute a SQL statement

    row = s3_cursor.fetchone()
    print(f"row = {row}")
    print('-' * 60)

    for row in s3_cursor.fetchmany(5):
        print(row)
    print('-' * 60)

    for row in s3_cursor.fetchall():
        print(row)
    print()
```

db_sqlite_fetch.py

```
row = (1, 'George', 'Washington', 'no party')
-----
(2, 'John', 'Adams', 'Federalist')
(3, 'Thomas', 'Jefferson', 'Democratic - Republican')
(4, 'James', 'Madison', 'Democratic - Republican')
(5, 'James', 'Monroe', 'Democratic - Republican')
(6, 'John Quincy', 'Adams', 'Democratic - Republican')
-----
(7, 'Andrew', 'Jackson', 'Democratic')
(8, 'Martin', 'Van Buren', 'Democratic')
```

...

```
(39, "James Earl 'Jimmy'", 'Carter', 'Democratic')
(40, 'Ronald Wilson', 'Reagan', 'Republican')
(41, 'George Herbert Walker', 'Bush', 'Republican')
(42, "William Jefferson 'Bill'", 'Clinton', 'Democratic')
(43, 'George Walker', 'Bush', 'Republican')
(44, 'Barack Hussein', 'Obama', 'Democratic')
(45, 'Donald John', 'Trump', 'Republican')
(46, 'Joseph Robinette', 'Biden', 'Democratic')
(47, 'Donald John', 'Trump', 'Republican')
```

Non-query statements

- Update database
- Returns count of rows affected
- Changes must be committed

The `execute()` method is also used to execute non-query statements, such as **CREATE**, **ALTER**, **UPDATE**, and **DROP**.

As with queries, the first argument is a string containing one SQL statement.

For most DB packages, `execute()` returns the number of rows affected.

To make changes to the database permanent, changes must be committed with `CONNECTION.commit()`.

Example

db_sqlite_add_row.py

```
from datetime import date
import sqlite3

with sqlite3.connect("../DATA/presidents.db") as s3conn: # connect to database

    sql_insert = """
        insert into presidents
        (lastname, firstname, termstart, termend, birthplace, birthstate, birthdate,
        deathdate, party)
        values ('Ramirez', 'Mary', '2025-01-20', null, 'Topeka',
        'Kansas', '1968-09-22', null, 'Independent')
    """

    cursor = s3conn.cursor()

    try:
        cursor.execute(sql_insert)
    except (sqlite3.OperationalError, sqlite3.DatabaseError, sqlite3.DataError) as err:
        print(err)
        s3conn.rollback()
    else:
        s3conn.commit()
    finally:
        cursor.close()
```

Example

db_sqlite_delete_row.py

```
from datetime import date
import sqlite3

with sqlite3.connect("../DATA/presidents.db") as conn: # connect to DB

    sql_delete = """
        delete from presidents
        where TERMNUM = 48
    """

    cursor = conn.cursor() # get a cursor

    try:
        cursor.execute(sql_delete)
    except (sqlite3.DatabaseError, sqlite3.OperationalError, sqlite3.DataError) as err:
        print(err)
        conn.rollback()
    else:
        conn.commit()
    finally:
        cursor.close()
```

SQL Injection

- Hijacks SQL code
- Result of string formatting
- Always use parameterized statements

One kind of vulnerability in SQL code is called *SQL injection*. This happens when using string formatting and raw user input to build SQL statements. An attacker can embed malicious SQL commands in input data.

Since the programmer is generating the SQL code as a string, there is no way to check for malicious SQL code. It is best practice to use parameterized statements, which prevents any user input from being *injected* into the SQL statement.



see <http://www.xkcd.com/327> for a well-known web comic on this subject.

Example

db_sql_injection.py

```
#  
good_input = 'Google'  
malicious_input = "'; drop table customers; -- " # input would come from a web form, for  
instance  
  
naive_format = "select * from customers where company_name = '{}' and company_id != 0"  
  
good_query = naive_format.format(good_input) # string formatting naively adds the user  
input to a field, expecting only a customer name  
malicious_query = naive_format.format(malicious_input) # string formatting naively adds  
the user input to a field, expecting only a customer name  
  
print("Good query:")  
print(good_query) # non-malicious input works fine  
print()  
  
print("Bad query:")  
print(malicious_query) # query now drops a table ('--' is SQL comment)
```

db_sql_injection.py

```
Good query:  
select * from customers where company_name = 'Google' and company_id != 0
```

```
Bad query:  
select * from customers where company_name = ''; drop table customers; -- ' and  
company_id != 0
```

Parameterized Statements

- Prevent SQL injection
- More efficient updates
- Use placeholders in query
 - Placeholders vary by DB
- Pass iterable of parameters
- Use cursor.execute() or cursor.executemany()

When using parameterized statements, you specify a SQL statement as usual, but use placeholders for the user-supplied data. Add an argument to `.execute()` containing the values for the placeholders. Placeholders can be either positional or named. Some DBMSs support both.

Parameterized statements not only protect against SQL injection, but they also are usually more efficient. The DBMS only parses the SQL statement once, then fills in data on each call to `.execute()`.

All SQL statements may be parameterized, including queries.

Different database modules use different placeholders. To see what kind of placeholder a module uses, check `_package_.paramstyle`.

Positional parameters

Positional parameters are filled in from an iterable. Values are filled in from left to right.

psycopg

```
cursor.execute("insert into users (name, id) values (%s, %s)", ["Bob", 123])
```

sqlite3

```
cursor.execute("insert into users (name, id) values (?, ?)", ["Bob", 123])
```

oracledb

```
cursor.execute("insert into users (name, id) values (:1, :2)", ["Bob", 123])
```

Named parameters

Named parameters are filled in from a dictionary. They have the advantage of not needing to be in a particular order.



oracledb supports specifying parameter values as named arguments to `.execute()` as well as a dictionary.

psycopg

```
cursor.execute(  
    "insert into users (name, id) values (%(name)s, %(id)s),  
    {'name':'Bob", 'id':123}  
)
```

sqlite3

```
cursor.execute(  
    "insert into users (name, id) values (:name, :id),  
    {'name':'Bob", 'id':123}  
)
```

oracledb

```
cursor.execute(  
    "insert into users (name, id) values (:name, :id),  
    name='Bob", id=123  
)  
  
cursor.execute(  
    "insert into users (name, id) values (:name, :id),  
    {'name':'Bob", 'id':123}  
)
```

executemany()

`executemany()` takes a non-query statement plus an iterable of parameter iterables (such as a list of tuples or a list of dictionaries). It will execute the query once for each item in the outer iterable.

`executemany()` may only be used for DML statements — another name for non-query statements.

```
users = [
    ('Bob', 123),
    ('Raul', 88),
    ('Lakshmi', 17),
]

cursor.executemany(
    "insert into users (name, id) values (%s, %s)",
    users
)
```

```
users = [
    {'name': 'Bob', 'id': 123},
    {'name': 'Raul', 'id': 88},
    {'name': 'Lakshmi', 'id': 17},
]

cursor.executemany(
    "insert into users (name, id) values (%(name)s, %(id)s)",
    users
)
```

Table 2. Placeholders for SQL Parameters

Python package	Positional Placeholder	Named Placeholder
pymysql	%s	%(param_name)s
oracledb	:1, :2, etc or :param_name	:param_name
pyodbc	?	None
pymssql	%s or %d	%(param_name)s <i>not</i> %(param_name)d
Psycopg	%s	%(param_name)s
SQLite	?	:param_name

Table 3. Placeholder types

Type	Description	Example
qmark	Question mark	WHERE name=?
numeric	Numeric, positional style	WHERE name=:1
named	Named style	WHERE name=:name
format	ANSI C printf format codes	WHERE name=%s
pyformat	Python extended format codes	WHERE name=%(name)s



package.paramstyle contains the parameter type for the given package.

Example

db_sqlite_parameterized.py

```
import sqlite3

# list of one-element tuples -- each tuple provides parameters for
# a SQL statement.
TERMS_TO_UPDATE = [(1,), (5,), (19,), (22,), (36,)]

PARTY_UPDATE = '''
update presidents
set party = "SURPRISE!"
where termnum = ?
''' # ? is SQLite3 placeholder for SQL statement parameter; different DBMSs use
      different placeholders

PARTY_QUERY = """
select termnum, firstname, lastname, party
from presidents
where termnum = ?
"""

with sqlite3.connect("../DATA/presidents.db") as s3conn:
    s3cursor = s3conn.cursor()
    # second argument to executemany() is iterable of parameters
    # each parameter is an iterable of values to fill in the placeholders
    s3cursor.executemany(PARTY_UPDATE, TERMS_TO_UPDATE)

    for param in TERMS_TO_UPDATE:
        term = param[0]
        s3cursor.execute(PARTY_QUERY, (term,))
        print(s3cursor.fetchone())
```

db_sqlite_parameterized.py

```
(1, 'George', 'Washington', 'SURPRISE!')
(5, 'James', 'Monroe', 'SURPRISE!')
(19, 'Rutherford Birchard', 'Hayes', 'SURPRISE!')
(22, 'Grover', 'Cleveland', 'SURPRISE!')
(36, 'Lyndon Baines', 'Johnson', 'SURPRISE!')
```

Example

db_sqlite_restore_parties.py

```
import sqlite3

# be sure to put values in correct order for placeholders
RESTORE_DATA = [
    ('no party', 1),
    ('Democratic - Republican', 5),
    ('Republican', 19),
    ('Democratic', 22),
    ('Democratic', 36),
]

PARTY_UPDATE = '''
update presidents
set party = ?
where termnum = ?

''' # ? is SQLite3 placeholder; other DBMSs may use other placeholders

PARTY_QUERY = """
select termnum, firstname, lastname, party
from presidents
where termnum = ?
"""

with sqlite3.connect("../DATA/presidents.db") as s3conn:
    s3cursor = s3conn.cursor()
    s3cursor.executemany(PARTY_UPDATE, RESTORE_DATA)
    s3conn.commit()

    # _ is throwaway variable
    for _, termnum in RESTORE_DATA:
        s3cursor.execute(PARTY_QUERY, [termnum])
        print(s3cursor.fetchone())
```

db_sqlite_restore_parties.py

```
(1, 'George', 'Washington', 'no party')
(5, 'James', 'Monroe', 'Democratic - Republican')
(19, 'Rutherford Birchard', 'Hayes', 'Republican')
(22, 'Grover', 'Cleveland', 'Democratic')
(36, 'Lyndon Baines', 'Johnson', 'Democratic')
```

Example

db_sqlite_bulk_insert.py

source

```
import sqlite3
import os
import csv

DATA_FILE = '../DATA/fruit_data.csv'

DB_NAME = 'fruits.db'
DB_TABLE = 'fruits'

SQL_CREATE_TABLE = f"""
create table {DB_TABLE} (
    id integer primary key,
    name varchar(30),
    unit varchar(30),
    unitprice decimal(6, 2)
)
""" # SQL statement to create table

SQL_INSERT_ROW = f'''
insert into {DB_TABLE} (name, unit, unitprice) values (?, ?, ?)
''' # parameterized SQL statement to insert one record

SQL_SELECT_ALL = f"""
select name, unit, unitprice from {DB_TABLE}
"""

def main():
    """
    Program entry point.

    :return: None
    """

    conn, cursor = get_connection()
    create_database(cursor)
    populate_database(conn, cursor)
    read_database(cursor)

    cursor.close()
    conn.close()
```

```
def get_connection():
    """
    Get a connection to the PRODUCE database

    :return: SQLite3 connection object.
    """
    if os.path.exists(DB_NAME):
        os.remove(DB_NAME) # remove database if it exists

    conn = sqlite3.connect(DB_NAME) # connect to (new) database
    cursor = conn.cursor()
    return conn, cursor

def create_database(cursor):
    """
    Create the fruit table

    :param conn: The database connection
    :return: None
    """
    cursor.execute(SQL_CREATE_TABLE) # run SQL to create table

def populate_database(conn, cursor):
    """
    Add rows to the fruit table

    :param conn: The database connection
    :return: None
    """
    with open(DATA_FILE) as file_in:
        fruit_data = csv.reader(file_in, quoting=csv.QUOTE_NONNUMERIC)

        for row in fruit_data:
            try:
                # add a row to the table
                cursor.execute(SQL_INSERT_ROW, row)
            except sqlite3.DatabaseError as err:
                print(err)
                conn.rollback()
            else:
                # commit the inserts; without this, no data would be saved
                conn.commit()

def read_database(cursor):
    cursor.execute(SQL_SELECT_ALL)
    for name, unit, unitprice in cursor.fetchall():
```

```
print(f'{name:12s} {unitprice:5.2f}/{unit}')
```

```
if __name__ == '__main__':
    main()
```

db_sqlite_bulk_insert.py

```
pomegranate 0.99/each
cherry        2.25/pound
apricot       3.49/pound
date          1.20/pound
apple         0.55/pound
lemon         0.69/each
kiwi          0.88/each
orange         0.49/each
lime           0.49/each
watermelon    4.50/each
guava          2.88/pound
papaya         1.79/pound
fig            2.29/pound
pear           1.10/pound
banana         0.65/pound
```

Example

db_sqlite_execute_many.py

```
import sqlite3
import os
import csv

DATA_FILE = '../DATA/fruit_data.csv'

DB_NAME = 'fruits.db'
DB_TABLE = 'fruits'

SQL_CREATE_TABLE = f"""
create table {DB_TABLE} (
    id integer primary key,
    name varchar(30),
    unit varchar(30),
    unitprice decimal(6, 2)
)
"""
# SQL statement to create table

SQL_INSERT_ROW = f"""
insert into {DB_TABLE} (name, unit, unitprice) values (?, ?, ?)
"""
# parameterized SQL statement to insert one record

SQL_SELECT_ALL = f"""
select name, unit, unitprice from {DB_TABLE}
"""

def main():
    """
    Program entry point.

    :return: None
    """

    conn, cursor = get_connection()
    create_database(cursor)
    populate_database(conn, cursor)

    # read database to confirm inserts
    read_database(cursor)

    cursor.close()
    conn.close()
```

```
def get_connection():
    """
    Get a connection to the PRODUCE database

    :return: SQLite3 connection object.
    """
    if os.path.exists(DB_NAME):
        os.remove(DB_NAME) # remove existing database if it exists

    conn = sqlite3.connect(DB_NAME) # connect to (new) database
    cursor = conn.cursor()
    return conn, cursor

def create_database(cursor):
    """
    Create the fruit table

    :param conn: The database connection
    :return: None
    """
    cursor.execute(SQL_CREATE_TABLE) # run SQL to create table

def populate_database(conn, cursor):
    """
    Add rows to the fruit table

    :param conn: The database connection
    :return: None
    """
    with open(DATA_FILE) as file_in:
        fruit_data = csv.reader(file_in, quoting=csv.QUOTE_NONNUMERIC)

        try:
            # iterate over rows of input
            # and add each row to database
            # fruit data is iterable of iterables
            # OR, iterable of dictionaries
            cursor.executemany(SQL_INSERT_ROW, fruit_data)
        except sqlite3.DatabaseError as err:
            print(err)
            conn.rollback()
        else:
            # Commit the inserts. Without this, no data
            # would be saved
            conn.commit()
```

```
def read_database(cursor):
    cursor.execute(SQL_SELECT_ALL)
    for name, unit, unitprice in cursor.fetchall():
        print(f'{name:12s} {unitprice:5.2f}/{unit}')

if __name__ == '__main__':
    main()
```

db_sqlite_execute_many.py

```
pomegranate 0.99/each
cherry        2.25/pound
apricot       3.49/pound
date          1.20/pound
apple         0.55/pound
lemon         0.69/each
kiwi          0.88/each
orange         0.49/each
lime           0.49/each
watermelon    4.50/each
guava          2.88/pound
papaya         1.79/pound
fig            2.29/pound
pear           1.10/pound
banana         0.65/pound
```

Metadata

- **cursor.description** returns tuple of tuples
- Fields
 - name
 - type_code
 - display_size
 - internal_size
 - precision
 - scale
 - null_ok

Once a query has been executed, the cursor's **description** attribute is a tuple with metadata about the columns in the query. It contains one tuple for each column in the query, containing 7 values describing the column.

For instance, to get the names of the columns, you could say

```
names = [d[0] for d in cursor.description]
```

For non-query statements, **CURSOR.description** returns **None**.

The names are based on the query (with possible aliases), and not necessarily on the names in the table.



Not all of the fields will necessarily be populated. For instance, **sqlite3** only provides column names.

Example

db_sqlite_metadata.py

```
"""
Provide metadata (tables and column names) for a Sqlite3 database
"""

from pprint import pprint
import sqlite3

DB_NAME = "../DATA/presidents.db"
TABLE_QUERY = '''select * from presidents where 1 = 2'''

def main():
    cursor = connect_to_db(DB_NAME)
    show_metadata(cursor)

def connect_to_db(database_file):
    with sqlite3.connect(database_file) as s3conn:
        return s3conn.cursor()

def show_metadata(cursor):
    cursor.execute(TABLE_QUERY)
    pprint(cursor.description)
    print()
    column_names = [column_data[0] for column_data in cursor.description]
    print(f"column_names = {column_names}")

if __name__ == '__main__':
    main()
```

db_sqlite_metadata.py

```
(('termnum', None, None, None, None, None, None),  
 ('lastname', None, None, None, None, None, None),  
 ('firstname', None, None, None, None, None, None),  
 ('termstart', None, None, None, None, None, None),  
 ('termend', None, None, None, None, None, None),  
 ('birthplace', None, None, None, None, None, None),  
 ('birthstate', None, None, None, None, None, None),  
 ('birthdate', None, None, None, None, None, None),  
 ('deathdate', None, None, None, None, None, None),  
 ('party', None, None, None, None, None, None))  
  
column_names = ['termnum', 'lastname', 'firstname', 'termstart', 'termend', 'birthplace',  
 'birthstate', 'birthdate', 'deathdate', 'party']
```

Dictionary Cursors

- Indexed by column name
- Not standardized in the DB API

Some DB packages provide dictionary cursors, which return a dictionary for each row, instead of a tuple. The keys are the names of the columns, so columns can be accessed by name rather than position.

Each package that provides a dictionary cursor has its own way of creating a dictionary cursor, although they all work the same way.



The `sqlite3` package provides a `Row` cursor, which can be indexed by position or by column name.

Table 4. Builtin Dictionary Cursors

Package	How to get a dictionary cursor
pymssql	<pre>conn = pymssql.connect (..., as_dict=True) dcur = conn.cursor() <i>all cursors will be dict cursors</i></pre>
psycopg ¹²	<pre>import psycopg.extras conn = psycopg.connect(...) dcur = conn.cursor(cursor_factory=psycopg.extras.DictCursor) <i>only this cursor will be a dict cursor</i></pre>
sqlite3 ¹	<pre>conn = sqlite3.connect (...) dcur = conn.cursor() dcur.row_factory = sqlite3.Row <i>only this cursor will be a dict cursor</i> conn = sqlite3.connect (...) conn.row_factory = sqlite3.Row dcur = conn.cursor() <i>all cursors will be dict cursors</i></pre>
pymysql ¹	<pre>import pymysql.cursors conn = pymysql.connect(...) dcur = conn.cursor(pymysql.cursors.DictCursor) <i>only this cursor will be a dict cursor</i> conn = pymysql.connect(..., cursorclass = pymysql.cursors.DictCursor) <i>all cursors will be dict cursors</i></pre>
oracledb	<i>Not available</i>
pyodbc	<i>Not available</i>
pgdb	<i>Not available</i>

¹ Cursor supports indexing by either key value (dict style) or integer position (list style), as well as iteration.² Also supports `RealDictCursor` which is an actual dictionary, and `NamedTupleCursor`, which is an actual `namedtuple`

Example

db_sqlite_dict_cursor.py

```
import sqlite3

s3conn = sqlite3.connect("../DATA/presidents.db")
# uncomment to make _all_ cursors dictionary cursors
# conn.row_factory = sqlite3.Row

NAME_QUERY = '''
    select firstname, lastname
    from presidents
    where termnum < 5
'''

cur = s3conn.cursor()

# select first name, last name from all presidents
cur.execute(NAME_QUERY)

for row in cur.fetchall():
    print(row)
print('-' * 50)

dict_cursor = s3conn.cursor() # get a normal SQLite3 cursor

# make _this_ cursor a dictionary cursor
dict_cursor.row_factory = sqlite3.Row # set the row factory to be a Row object

# Row objects are dict/list hybrids -- row[name] or row[pos]

# select first name, last name from all presidents
dict_cursor.execute(NAME_QUERY)

for row in dict_cursor:
    print(row['firstname'], row['lastname']) # index row by column name

print('-' * 50)
```

db_sqlite_dict_cursor.py

```
('George', 'Washington')
('John', 'Adams')
('Thomas', 'Jefferson')
('James', 'Madison')
```

```
-----  
George Washington  
John Adams  
Thomas Jefferson  
James Madison  
-----
```

Generic alternate cursors

- Create generator function
 - Get column names from `cursor.description()`
 - For each row
 - Make object from column names and values
 - Dictionary
 - Named tuple
 - Dataclass

For database modules that don't provide a dictionary cursor, the `iterrows_asdict()` function described below can be used with a cursor from any DB API-compliant package.

The example uses the metadata from the cursor to get the column names, and forms a dictionary by zipping the column names with the column values. `db_iterrows` also provides `iterrows_asnamedtuple()`, which returns each row as a named tuple, and `iterrows_asdataclass()`, which returns each row as an instance of a custom dataclass.

The functions in `db_iterrows` return generator objects. When you loop over the generator object, each element is a dictionary, named tuple, or instance of a dataclass, depending on which function you called.

Example

db_iterrows.py

```
"""
Generic functions that can be used with any DB API compliant
package.

To use, pass in a cursor after execute()-ing a
SQL query. Then iterate over the generator that is
returned
"""

from collections import namedtuple
from dataclasses import make_dataclass

def get_column_names(cursor):
    return [desc[0] for desc in cursor.description]

def iterrows_asdict(cursor):
    '''Generate rows as dictionaries'''
    column_names = get_column_names(cursor)
    for row in cursor.fetchall():
        row_dict = dict(zip(column_names, row))
        yield row_dict

def iterrows_asnamedtuple(cursor):
    '''Generate rows as named tuples'''
    column_names = get_column_names(cursor)
    Row = namedtuple('Row', column_names)
    for row in cursor.fetchall():
        yield Row(*row)

def iterrows_asdataclass(cursor):
    '''Generate rows as dataclass instances'''
    column_names = get_column_names(cursor)
    Row = make_dataclass('row_tuple', column_names)

    for row in cursor.fetchall():
        yield Row(*row)
```

Example

db_sqlite_iterrows.py

```
"""
Generic functions that can be used with any DB API compliant
package.

To use, pass in a cursor after execute()-ing a
SQL query. Then iterate over the generator that is
returned
"""

import sqlite3
from db_iterrows import *

sql_select = """
SELECT firstname, lastname, party
FROM presidents
WHERE termnum > 39
"""

conn = sqlite3.connect("../DATA/presidents.db")

cursor = conn.cursor()

cursor.execute(sql_select)

for row in iterrows_asdict(cursor):
    print(row['firstname'], row['lastname'], row['party'])

print('-' * 60)

cursor.execute(sql_select)

for row in iterrows_asnamedtuple(cursor):
    print(row.firstname, row.lastname, row.party)

print('-' * 60)

cursor.execute(sql_select)

for row in iterrows_asdataclass(cursor):
    print(row.firstname, row.lastname, row.party)
```

db_sqlite_iterrows.py

```
Ronald Wilson Reagan Republican
George Herbert Walker Bush Republican
William Jefferson 'Bill' Clinton Democratic
George Walker Bush Republican
Barack Hussein Obama Democratic
Donald John Trump Republican
Joseph Robinette Biden Democratic
Donald John Trump Republican
```

```
Ronald Wilson Reagan Republican
George Herbert Walker Bush Republican
William Jefferson 'Bill' Clinton Democratic
George Walker Bush Republican
Barack Hussein Obama Democratic
Donald John Trump Republican
Joseph Robinette Biden Democratic
Donald John Trump Republican
```

```
Ronald Wilson Reagan Republican
George Herbert Walker Bush Republican
William Jefferson 'Bill' Clinton Democratic
George Walker Bush Republican
Barack Hussein Obama Democratic
Donald John Trump Republican
Joseph Robinette Biden Democratic
Donald John Trump Republican
```

Transactions

- Transactions allow safer control of updates
- `commit()` to make database changes permanent
- `rollback()` to discard changes

Sometimes a database task involves more than one change to your database (i.e., more than one SQL statement). You don't want the first SQL statement to succeed and the second to fail; this would leave your database in a corrupt state.

To be certain of data integrity, use **transactions**. This lets you make multiple changes to your database and only commit the changes if all the SQL statements were successful.

For all packages using the Python DB API, a transaction is started when you connect. At any point, you can call `CONNECTION.commit()` to save the changes, or `CONNECTION.rollback()` to discard the changes. If you don't call `commit()` after modifying a table, the data will not be saved.

Example

```
for values in list_of_tuples:  
    try:  
        cursor.execute(query,values)  
    except SQLError:  
        dbconn.rollback()  
    else:  
        dbconn.commit()
```

Autocommit

- Commits after every statement
- Not a best practice

Autocommit mode calls `commit()` after every non-query statement. This is not generally considered a best practice. See the table below for how autocommit is implemented in various DB packages.

Table 5. How to turn on autocommit

Package	Method/Attribute
oracledb	<code>conn.autocommit = True</code>
ibm_db_api	<code>conn.set_autocommit(True)</code>
pymysql	<code>pymysql.connect(…, autocommit=True)</code> Or <code>conn.autocommit(True)</code>
psycopg	<code>conn.autocommit = True</code>
sqlite3	<code>sqlite3.connect(dbname, isolation_level=None)</code>



pymysql only supports transaction processing when using the **InnoDB** engine

Object-relational Mappers

- No SQL required
- Maps a class to a table
- All DB work is done by manipulating objects
- Most popular Python ORMs
 - SQLAlchemy
 - Django (which is a complete web framework)

An Object-relational mapper is a module or framework that creates a level of abstraction above the actual database tables and SQL queries. As the name implies, a Python class (object) is mapped to the actual table.

The two most popular Python ORMs are SQLAlchemy which is a standalone ORM, and Django ORM. Django is a comprehensive Web development framework, which provides an ORM as a subpackage. SQLAlchemy is the most fully developed package, and is the ORM used by Flask and some other Web development frameworks.

Instead of querying the database, you call a search method on an object representing a table. To add a row to the table, you create a new instance of the table class, populate it, and call a method like `save()`. You can create a large, complex database system, complete with foreign keys, composite indices, and all the other attributes near and dear to a DBA, without writing the first line of SQL.

You can use Python ORMs in two ways.

One way is to design the database with the ORM. To do this, you create a class for each table in the database, specifying the columns with predefined classes from the ORM. Then you run an ORM command which executes the queries needed to build the database. If you need to make changes, you update the class definitions, and run an ORM command to synchronize the actual DBMS to your classes.

The second way is to map tables to an existing database. You create the classes to match the schemas that have already been defined in the database. Both SQLAlchemy and the Django ORM have tools to automate this process.

NoSQL

- Non-relational database
- Document-oriented
- Can be hierarchical (nested)
- Examples
 - MongoDB
 - Cassandra
 - Redis

A current trend in data storage are called "NoSQL" or non-relational databases. These databases consist of *documents*, which are indexed, and may contain nested data.

NoSQL databases don't contain tables, and do not have relations.

While relational databases are great for tabular data, they are not as good a fit for nested data. Geospatial, engineering diagrams, and molecular modeling can have very complex structures. It is possible to shoehorn such data into a relational database, but a NoSQL database might work much better. Another advantage of NoSQL is that it can adapt to changing data structures, without having to rebuild tables if columns are added, deleted, or modified.

Some of the most common NoSQL database systems are MongoDB, Cassandra and Redis.

Chapter 1 Exercises

Exercise 1-1

Part A (president_sqlite.py)

For this exercise, use the SQLite database named `presidents.db` in the DATA folder. It has the following layout

Table 6. Layout of President Table

Field Name	SQLite Data Type	Python Data type	Null	Default
termnum	int(11)	int	YES	NULL
lastname	varchar(32)	str	YES	NULL
firstname	varchar(64)	str	YES	NULL
termstart	date	date	YES	NULL
termend	date	date	YES	NULL
birthplace	varchar(128)	str	YES	NULL
birthstate	varchar(32)	str	YES	NULL
birthdate	date	date	YES	NULL
deathdate	date	date	YES	NULL
party	varchar(32)	str	YES	NULL

Copy the `president.py` module in the top level folder of the student files to `president_sqlite.py` and then modify `president_sqlite.py` to get its data from the Presidents table, as described above, rather than from `presidents.txt`.

Part B (president_main.py)

Copy `president_main.py` to `president_main_sqlite.py` and modify it to import the `President` class from the `president_sqlite` module. It should import the `President` class from this new module that uses the database instead of a text file, but otherwise work the same as before.

Exercise 1-2 (add_pres_sqlite.py)

Add another president to the presidents database. Just make up the data for your new president.

SQL syntax for adding a record is

```
INSERT INTO table ("COL1-NAME",...) VALUES ("VALUE1",...)
```

To do a parameterized insert (the right way!):

```
INSERT INTO table ("COL1-NAME",...) VALUES (?,?,...)
```

Index

@

□0□, 4, 8, 14, 21, 31

A

API, 2

C

Cassandra, 46

commit, 42

connection object, 7

context manager, 4

cursor, 7

cursor object, 7

cx_oracle, 3

D

database programming, 2

database server, 4

Databricks, 3

DB API, 2

dictionary cursor

emulating, 39

dictionary cursors, 34

Django, 45

Django ORM, 45

DML, 21

E

Elasticsearch, 3

Excel, 3

executing SQL statements, 11

F

Firebird (and Interbase, 3

G

Google BigQuery, 3

I

IBM DB2, 3

ibm-db, 3

Informix, 3

informixdb, 3

ingmod, 3

Ingres, 3

K

KInterbasDB, 3

M

metadata, 31

Microsoft SQL Server, 3

MongoDB, 46

MySQL, 3

N

namedtuple cursor, 39

non-query statement, 21

non-relational, 46

NoSQL, 46

O

Object-relational mapper, 45

ODBC, 3

Oracle, 3

oracledb, 3

ORM, 45

P

parameterized SQL statements, 21

PostgreSQL, 3

psycopg, 3

pymssql, 3

pymysql, 3

pyodbc, 3

R

Redis, 46

rollback, 42

S

SAP DB, 3

sapdbapi, 3

SQL code, 7

SQL data integrity, [42](#)

SQL injection, [17](#)

SQL queries, [11](#)

SQLAlchemy, [45](#)

SQLite, [3](#)

sqlite3, [3](#)

Sybase, [3, 3](#)

T

transactions, [42](#)